

Model Optimization and Tuning Phase Report

Date	23 September 2024
Team ID	LTVIP2024TMID24997
Project Title	SmartLender - Applicant Credibility Prediction for Loan Approval
Maximum Marks	10 Marks

Model Optimization and Tuning Phase

The Model Optimization and Tuning Phase involves refining machine learning models for peak performance. It includes optimized model code, fine-tuning hyperparameters, comparing performance metrics, and justifying the final model selection for enhanced predictive accuracy and efficiency.

Hyperparameter Tuning Documentation (6 Marks):

Model	Tuned Hyperparameters	Optimal Values
Linear Regression	<pre>import numpy as np from sklearn.linear_model import LinearRegression from sklearn.metrics import mean_squared_error, r2_score # Model training linear_reg = LinearRegression() linear_reg.fit(x_train, y_train) # Prediction linear_pred = linear_reg.predict(x_test) # Evaluation print("MSE:", mean_squared_error(y_test, linear_pred)) print("R2 Score:", r2_score(y_test, linear_pred)) # Custom accuracy score based on a tolerance level def accuracy_score(y_true, y_pred, tolerance=0.1): return np.mean(np.abs(y_true - y_pred) <= tolerance)</pre>	<pre># Evaluation print("MSE:", mean_squared_error(y_test, linear_pred)) print("R2 Score:", r2_score(y_test, linear_pred)) # Calculate accuracy using R2 score accuracy = r2_score(y_test, linear_pred) print("R2 Score (Accuracy):", accuracy)</pre> <p>MSE: 0.011427751339300353 R² Score: 0.999997047756003 R² Score (Accuracy): 0.999997047756003</p>
Ridge regression	<pre>def accuracy_score(y_true, y_pred, tolerance=0.1): return np.mean(np.abs(y_true - y_pred) <= tolerance) # Standardize features scaler = StandardScaler() # Create a pipeline with scaling and Ridge regression ridge_pipeline = Pipeline([('scaler', scaler), ('ridge', Ridge())]) # Hyperparameter tuning using GridSearchCV param_grid = { 'ridge__alpha': [0.01, 0.1, 1.0, 10.0, 100.0] # Range of alpha values } grid_search = GridSearchCV(ridge_pipeline, param_grid, cv=5, scoring='r2') grid_search.fit(x_train, y_train)</pre>	<pre># Evaluation print("MSE:", mean_squared_error(y_test, ridge_pred)) print("R2 Score:", r2_score(y_test, ridge_pred)) # Calculate accuracy using R2 score accuracy = r2_score(y_test, ridge_pred) print("R2 Score (Accuracy):", accuracy)</pre> <p>MSE: 0.0720158246881738 R² Score: 0.9999981395439939 R² Score (Accuracy): 0.9999981395439939</p>

Lasso regression	<pre># Define a custom accuracy score based on a tolerance level def accuracy_score(y_true, y_pred, tolerance=0.1): return np.mean(np.abs(y_true - y_pred) <= tolerance) # Standardize features scaler = StandardScaler() # Create a pipeline with scaling and lasso regression lasso_pipeline = Pipeline([('scaler', scaler), ('lasso', Lasso())]) # Hyperparameter tuning using GridSearchCV param_grid = { 'lasso_alpha': [0.01, 0.1, 1.0, 10.0, 100.0] # Range of alpha values to test } grid_search = GridSearchCV(lasso_pipeline, param_grid, cv=5, scoring='r2') # Cross-validation grid_search.fit(x_train, y_train)</pre>	<pre># Calculate accuracy using R² score accuracy = r2_score(y_test, lasso_pred) print("R² Score (Accuracy):", accuracy) print("Best Parameters:", grid_search.best_params_) MSE: 0.012321835829863403 R² Score: 0.9999996816778316 R² Score (Accuracy): 0.9999996816778316 Best Parameters: {'lasso_alpha': 0.01}</pre>
Decision Tree	<pre># Define hyperparameter grid for Decision Tree Regression dt_param_dist = { 'dt__max_depth': [None, 3, 5, 10, 15], # Maximum depth 'dt__min_samples_split': [2, 5, 10], # Minimum samples to split 'dt__min_samples_leaf': [1, 2, 4, 6], # Minimum samples in leaf 'dt__max_features': ['auto', 'sqrt', 'log2'] # Number of features to consider } # Use RandomizedSearchCV for hyperparameter tuning dt_search = RandomizedSearchCV(dt_pipeline, dt_param_dist, n_iter=50, cv=5, scoring='r2', n_jobs=-1, random_state=42) # Fit the model dt_search.fit(X_train, y_train)</pre>	<pre>best_dt = grid_search_dt.best_estimator_ print("Best Hyperparameters for Decision Tree:", grid_search_dt.best_params_) y_pred_dt = best_dt.predict(X_test) print("Decision Tree Accuracy: {:.2f}".format(accuracy_score(y_test, y_pred_dt))) Fitting 5 folds for each of 216 candidates, totalling 1080 fits Best Hyperparameters for Decision Tree: {'criterion': 'gini', 'max_depth': 10} Decision Tree Accuracy: 0.74</pre>

Performance Metrics Comparison Report (2 Marks):

Model	Baseline Metric	Optimized Metric
Linear Regression	R2: 0.999	R2: 0.999
Ridge Regression	R2: 0.998	R2: 0.998
Lasso Regression	R2: 0.996	R2: 0.996
Decision Tree Regressor	R2: 0.7445	R2: 0.7445

Final Model Selection Justification (2 Marks):

Final Model	Reasoning
Linear Regression	The Linear Regression model was chosen as the final optimized model because it exhibited the highest R-squared value (0.999), indicating a strong fit to the data. Additionally, it had a lower MSE (0.0114) and Accuracy (99%) compared to other models, suggesting superior predictive accuracy.

